

```

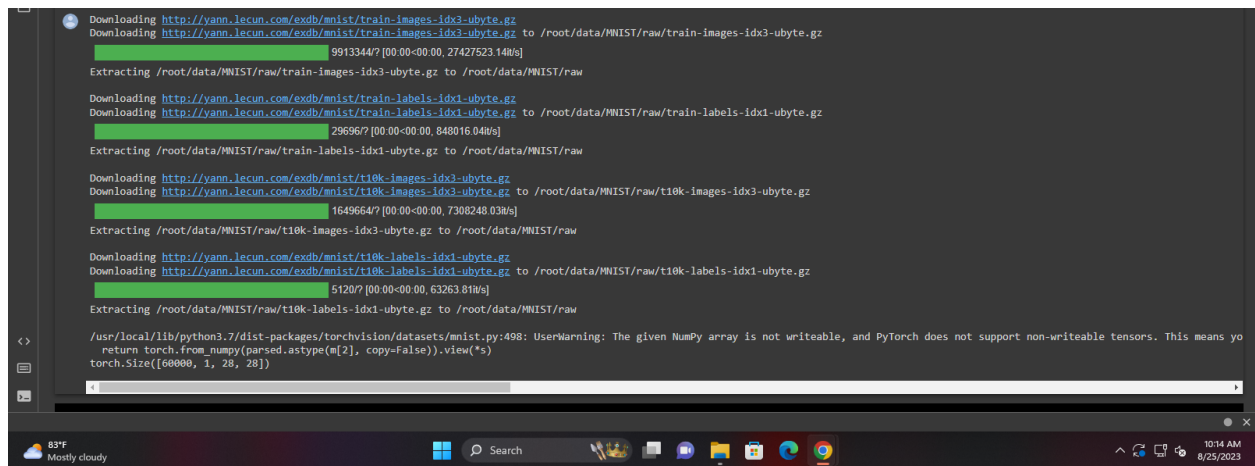
import torch
import torchvision
from torch import nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
rng = np.random.default_rng(123456)

```

```

data = torchvision.datasets.MNIST(root='~/data', download=True)
data = data.data
data = data.float() / 255.
data = data.view(-1, 1, 28, 28)
print(data.shape)

```



```

class AutoEncoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Flatten(),
            nn.Linear(28*28, 100),
            nn.ReLU(),
            nn.Linear(100, 10),
            nn.ReLU(),
        )
        self.decoder = nn.Sequential(
            nn.Linear(10, 100),
            nn.ReLU(),
        )

```

```

        nn.Linear(100, 28*28),
        nn.Sigmoid()
    )

    def encode(self, x):
        return self.encoder(x)

    def decode(self, x):
        x = self.decoder(x)
        return x.view(-1, 1, 28, 28)

    def forward(self, x):
        return self.decode(self.encode(x))

```

```

model = AutoEncoder().cuda()
opt = torch.optim.Adam(model.parameters())

```

```

for epoch in range(25):
    print(f'Epoch {epoch+1}/25')
    for i in range(0, data.shape[0], 32):
        x = data[i:i+32].cuda()
        x_rec = model(x)
        loss = F.binary_cross_entropy(x_rec, x)

        opt.zero_grad()
        loss.backward()
        opt.step()

    data = data[rng.permutation(len(data))]
    print(f'\tloss: {loss.item():.4f}')

```

The screenshot shows a Google Colab notebook titled 'autoencoder_example.ipynb'. The code cell contains a loop that prints the loss for each epoch from 1 to 19. The loss values are as follows:

Epoch	Loss
1/25	0.1552
2/25	0.1270
3/25	0.1206
4/25	0.1324
5/25	0.1237
6/25	0.1256
7/25	0.1357
8/25	0.1197
9/25	0.1368
10/25	0.1163
11/25	0.1418
12/25	0.1164
13/25	0.1288
14/25	0.1217
15/25	0.1252
16/25	0.1132
17/25	0.1180
18/25	0.1194
19/25	


The screenshot shows the continuation of the Google Colab notebook, displaying epochs 18 to 25. The loss values are as follows:

Epoch	Loss
18/25	0.1194
19/25	0.1155
20/25	0.1237
21/25	0.1141
22/25	0.1259
23/25	0.1217
24/25	0.1208
25/25	0.1151

```
plt.figure(figsize=(5,10))

for i in range(5):
    plt.subplot(5, 2, i*2+1, title=f'Train image')
    plt.imshow(np.squeeze(x[i].cpu()), cmap='gray')
    plt.axis('off')

    plt.subplot(5, 2, i*2+2, title='Reconstruction')
    with torch.no_grad(): plt.imshow(np.squeeze(x_rec[i].cpu()),
cmap='gray')
    plt.axis('off')
```



```
# Sample two random images and encode
f = model.encode(x[0:2])
f1, f2 = f[0].unsqueeze(0), f[1].unsqueeze(0)

# Show reconstructions of interpolated codes
plt.figure(figsize=(20,5))
reconstructions = []
for i in range(20):
    v = i/19.
    f_interp = f1*(1-v) + f2*v
    with torch.no_grad():
        x_rec_interp = np.squeeze(model.decode(f_interp).cpu())
        reconstructions.append(x_rec_interp)


plt.subplot(2,10,i+1)
plt.imshow(x_rec_interp, cmap='gray')
plt.axis('off')
```



```
from IPython.display import HTML
from matplotlib import animation

fig = plt.figure()
plt.axis('off')
artists = [[plt.imshow(img, animated=False, cmap='gray')] for
img in reconstructions]
ani = animation.ArtistAnimation(fig, artists, interval=200,
blit=False, repeat_delay=1000)

print('Latent interpolation')
HTML(ani.to_html5_video())
```



```
img1,img2 = x[0],x[2]
images = []
for i in range(20):
    v = i/19.
    img_interp = img1*(1-v) + img2*v
    images.append(np.squeeze(img_interp.cpu()))

fig = plt.figure()
plt.axis('off')
artists = [[plt.imshow(img, animated=False, cmap='gray')] for
img in images]
ani = animation.ArtistAnimation(fig, artists, interval=200,
blit=False, repeat_delay=1000)

print('Pixel interpolation')
HTML(ani.to_html5_video())
```

